
CSE 431 Computer Architecture Fall 2015

Chapter 4E: Static SuperScalar (SS) Datapaths

Mary Jane Irwin (www.cse.psu.edu/~mji)

[Adapted from *Computer Organization and Design, 5th Edition*,
Patterson & Hennessy, © 2014, MK
With additional thanks/credits to Amir Roth, Milo Martin, CIS/UPenn]

Reminders

□ This week

- Static SuperScalar processors – P&H 4.10
 - VLIW, in-order, and FGMT

□ Next week

- Dynamic SuperScalar processors – P&H 4.10
 - out-of-order and SMT (aka hyperthreading)

□ Reminders

- HW4 is due Oct 26th
- Quiz4 will close on Oct 23rd
- Second (and last) evening exam scheduled
 - Tuesday, **November 17**, 20:15 to 22:15, Location 22 Deike
 - Please let me know ASAP if you have a conflict !!

Review: Multiple-Issue Datapath Responsibilities

- ❑ Must handle, with a combination of hardware and software fixes, the fundamental limitations of
 - How many instructions to **issue** (send for execution) in one clock cycle
 - Storage (data) dependencies → **data hazards**
 - Limitation more severe in a in-order SuperScalar/VLIW processor due to (usually) low ILP
 - Procedural dependencies → **control hazards**
 - Ditto, but even more severe
 - Use dynamic branch prediction to help resolve the ILP issue
 - Use loop unrolling (in the compiler) to increase ILP
 - Resource conflicts → **structural hazards**
 - A multiple-issue datapath has a much larger number of potential resource conflicts
 - Functional units may have to arbitrate for result buses and RF write ports
 - Resource conflicts can be reduced by duplicating the resource or by pipelining the resource


Review: Overview of Dependence Analysis

- ❑ To what extent can the compiler (or the datapath) reorder instructions? Are there execution-order constraints?

original	possible?	possible?
instr 1 instr 2 consecutive	instr 2 instr 1 consecutive	instr 1 and instr 2 simultaneous

- ❑ **Instruction dependencies** imply that reordering instructions is not possible
 - true dependence (or, data dep., flow dep.) (cannot reorder)
 - $a = .$
 - $. = a$ **RAW, read after write**
 - anti-dependence (renaming allows reordering)
 - $. = a$
 - $a = .$ **WAR, write after read**
 - output dependence (renaming allows reordering)
 - $a = .$
 - $a = .$ **WAW, write after write**

Static SS IF Stage Challenges



- ❑ Wide instruction fetch: Fetching a 8B to 32B (2 to 8 instr's assuming 32b (4B) instr's) from the IM at once
 - Have to design the IM (I\$) to support wide fetch in one cycle
- ❑ How many **branches** do we allow in a fetch bundle?

Answer is usually only one (so that we only have to build one branch predictor).


 - Discard post-branch instr's in the fetch bundle if the prediction is "taken" which lowers the effective fetch width and the IPC
 - As we have seen, the compiler can help reduce the branch frequency with loop unrolling

CmpEn431 Chapter 4E.7

Irwin, PSU, 2015

Instruction Fetch Sequences


- ❑ Instruction run – number of (sequential) instructions (run length) fetched between **taken** branches
 - Instruction fetcher operates most efficiently when processing long runs – unfortunately runs are usually quite short (about six instr's)
- ❑ Example: for a 4-way fetcher, (instr fetch bandwidth of 4 instr's per cycle with branch prediction)
 - 8 instructions in 4 cycles – so a actual rate of only 2 instr's/cycle
- ❑ Fetcher can merge instr's from different runs, if it has a fetch rate faster than the decode rate, $\text{fetch: speed} < \text{ratio} >$
 - 8 instructions in 3 cycles



CmpEn431 Chapter 4E.8

Irwin, PSU, 2015


Speedups of Fetch Alternatives



□ A 4-way instr fetcher out performs a 2-way instr fetcher

- It has twice the potential instruction bandwidth
- But it requires twice as much decoder hardware to keep up

Static SS Dec Stage Challenges



□ Have to decode 2 to 8 instr's *at once* and decide which can issue (be sent to the Exec stage) *in parallel*

- Duplicated decoders
- Logic to determine if there are structural hazards and/or data dependencies in the current instr bundle or load-use hazards with the previous instr bundle
- Logic to **stall** conflicted instr's (and instr's in Fetch) for a cycle


□ Multiported RF – 4 read ports/2 write ports (2 instr's) up to 16 read ports/8 write ports (8 instr's)

- Larger area, latency, power, ...

Aside: A 2 Read Port, 1 Write Port RF

N Write ports: area, latency $\sim N^2$

2N Read ports: area, latency $\sim (2N)^2$



CmpEn431 Chapter 4E.11

Irwin, PSU, 2015

Dependency Checking


- ❑ Need to check for structural hazards (do the 2 (or 4, or 8) instr's need same FU's in EX ?)
 - If so need to either duplicate the FU's or stall one (or more) of the instr's in the bundle.
- ❑ Need to cross check for load-use hazards of the instr's in ID (the "use" instr's - for both of their **src** operands) to the instr's in EX (the "load" instr's). We have forwarding logic that can take care of all other inter-bundle RAW data hazards for our 5-stage pipeline.
- ❑ And need to check for **dst-src** (RAW) and **dst-dst** (WAW) dependencies between the instr's in the **same** instruction bundle in ID (intra-bundle RAW and WAW)
- ❑ Don't really have to check for WAR intra-bundle dependencies, why not?

CmpEn431 Chapter 4E.12


Irwin, PSU, 2015

2-way Dependency Checking

- Cross check for load-use hazards of the 2 instr's in ID (for both **src's**) to the 2 instr's in EX which gives 8 load-use dependency checks



- And check for 2 **dst-src** (RAW) and 1 **dst-dst** (WAW) dependencies between the 2 instr's in the **same** instr bundle in ID




Load-Use Stalls


- Cycle 1: Load-use hazard detected



- Cycle 2: Load-use hazard bubble insertion (instructions in FETCH stage and ID stage are stalled for one cycle)




- Cycle 3: MEM/WB to EX data forwarding will now happen



RAW (and WAW) Stalls


- ❑ Cycle 1: RAW hazard detected




- ❑ Cycle 2: W instr moves to EX stage, R instr in ID and instructions in FETCH stalled for one cycle



- ❑ Cycle 3: EX/MEM to EX data forwarding will now happen




Static SS Exec Stage Challenges




- ❑ Need multiple execution units, do we need N (N = # instr in an instr bundle) of every kind?
 - ALUs? FP dividers?
 - How many branches per bundle? (we already decided only one)
 - How many loads and/or stores per bundle?
- ❑ Usually some mix proportional to the instr mix
 - 2-way: 1 integer (branch, load, store, int) + 1 ALU (int, fp)
 - 4-way: 2 integer + 2 ALU

Static SS Mem Stage Challenges



- ❑ What about multiple loads and/or stores per cycle?
 - Probably only needed in 4-wide or greater
 - More important to support multiple loads than multiple stores
 - Instr mix: loads (~20% to 25%), stores (~10% to 15%)
- ❑ Have to design the DM (D\$) to support multiple loads/stores in one cycle (have assumed only one DM port to this point)
 - Multi-porting is expensive in terms of latency, area, and power
 - Banked (interleaved) memories

Static SS WB Stage Challenges



- ❑ For an N-wide machine, need 2N RF read ports and N write ports
 - Read ports: area, latency $\sim (2N)^2$
 - Write ports: area, latency $\sim N^2$
- ❑ May not use the max number of read and write ports
 - Read ports: not all instr's use two source operands; forwarding supplies many of the read values (but don't know that at RF read time, so it doesn't help reduce read port count)
 - Write ports: stores, branches (~35%) don't write to the RF

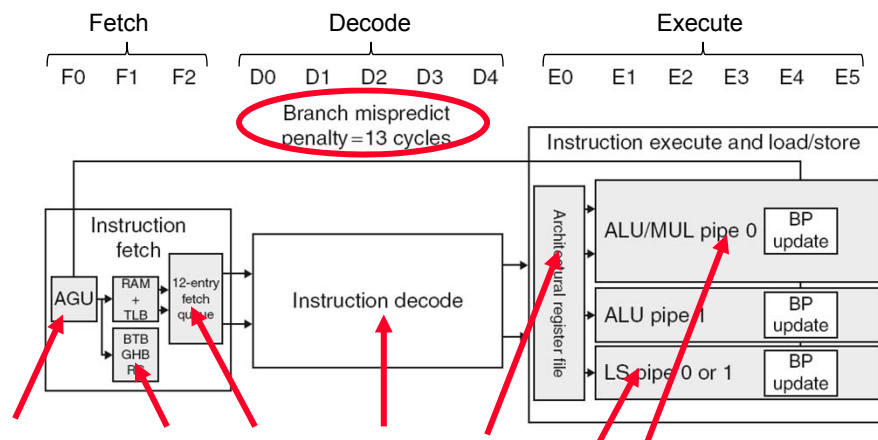
Trends in Static SS Datapath Design

	Pentium	PentiumII	Pentium4	Itanium	ItaniumII	Core2
Year	1993	1998	2001	2002	2004	2006
Width	2	3	3	3	6	4

- ❑ Issue width has saturated at 4- to 6-way for high-performance cores
 - The canceled DEC Alpha 21464 was an 8-way issue
 - Hardware or compiler “scheduling” needed to exploit 4- to 6-way effectively
 - VLIW or EPIC (Itanium)
- ❑ For good-performance, low-power cores, issue width is ~2
 - So advanced scheduling techniques not needed
 - Use multi-threading (stay tuned ...) to help cope with load-use hazards and cache misses


ARM Cortex A8 Pipeline

- ❑ 2-wide static (in-order) superscalar, 14-stage pipeline, 1GHz clock



ARM Cortex A8 Performance

- Ideal CPI is 0.5. For the median case (*gcc*), 80% of the stalls are due to pipeline hazards, 20% to memory stalls
 - Pipeline hazards are from branch mispredictions, structural hazards, and data dependencies between pairs
 - The compiler is the only thing that can help with structural hazards and data dependencies



Aside: CISC vs RISC vs Static SS vs VLIW

	CISC	RISC	Static Superscalar	VLIW
Instr size	variable size	fixed size	fixed size	fixed size (but large)
Instr format	variable format	fixed format	fixed format	fixed format
Registers	few, some special Limited # of ports	Many GP Limited # of ports	Many (more) GP Many ports	Many, many GP Many ports
Memory reference	embedded in many instr's	load/store	load/store	load/store
Key Issues	decode complexity	data forwarding, hazards	hardware instr dependency checks, data forwarding	(compiler) code scheduling


Multi-threading (MT)

- ❑ Even moderate static superscalars (e.g., 4-way) are not fully utilized
 - Average sustained IPC: 1.5–2 → < 50% utilization due to
 - Mispredicted branches
 - Cache misses, especially L1
 - Data dependences (e.g., load-use data hazards)
- ❑ Multi-threading (MT) to the rescue
 - Improve *utilization* of datapath components by multiplexing multiple (process) threads on single datapath
 - If one thread cannot fully utilize the datapath, maybe 2 or 4 (or 100) can

Multithreading Example


- ❑ Time evolution of issue slot

- 4-way datapath



Static SS

- ❑ # cycles? # wasted cycle slots?



Multithreaded Static SS

- ❑ Fill in with instructions from other threads – in this example we have 2 threads and change threads every cycle


- Completely removes load-use hazard empty slots
- Takes longer for the “red” thread to finish
 - With more threads, would take even longer
- Still have some noop slots (so wasted performance – stay tuned)

Alternative Multithreaded Implementations

- ❑ MT trades (single-thread) latency for throughput
 - Sharing the datapath degrades the **latency** of individual threads, but improves the aggregate latency of both threads
 - And it improves **utilization** of the datapath hardware
- ❑ Main questions: **thread scheduling policy** and **pipeline partitioning**
 - When to switch from one thread to another?
 - How exactly do threads share the pipelined datapath itself?
- ❑ Choices depends on what kind of latencies you want to tolerate and how much single thread performance you are willing to sacrifice
 - Coarse-grain multithreading (**CGMT**)
 - Fine-grain multithreading (**FGMT**) ←
 - Simultaneous multithreading (**SMT**)

Fine-Grain MultiThreading (FGMT)

- + Tolerates latencies (e.g., load-use hazards, L1 misses, mispredicted branches, etc.)
- Sacrifices significant single thread performance
- Need a *lot* of threads to get returns
- ❑ Thread scheduling policy
 - Switch threads every cycle (e.g., round-robin, can skip stalled threads)
- ❑ Pipeline partitioning
 - Dynamic, *no* pipeline flushing between threads



FGMT

Examples of FGMT Processors

❑ Sun's UltraSPARC T1 (Niagara)

- Many threads → many RF
- Power efficient

http://en.wikipedia.org/wiki/UltraSPARC_T1

Extreme examples

❑ Denelcor HEP (Burton Smith architect)

- So many threads (100+), it didn't even need caches
- Targeted for DoD, not successful commercially

http://en.wikipedia.org/wiki/Heterogeneous_Element_Processor

❑ CRAY MTA Threadstorm (another of Smith's designs)

- Again many threads (128) per core (8,000+) each with small thread state
- Thread-level context switch every instruction cycle

<http://www.jp.cray.com/downloads/XMT-Presentation.pdf>

FGMT Sharing Implementations Issues

❑ How do multiple threads share a single datapath?

- Different sharing mechanisms for different kinds of structures depending on what kind of **state** the structure stores

❑ **No state**: ALUs

- So can be dynamically shared

❑ **Persistent hard state (aka thread "context")**: PC, RFile

- So must be **replicated**

❑ **Persistent soft state**: caches, TLBs, bpred (BTB, BHT)

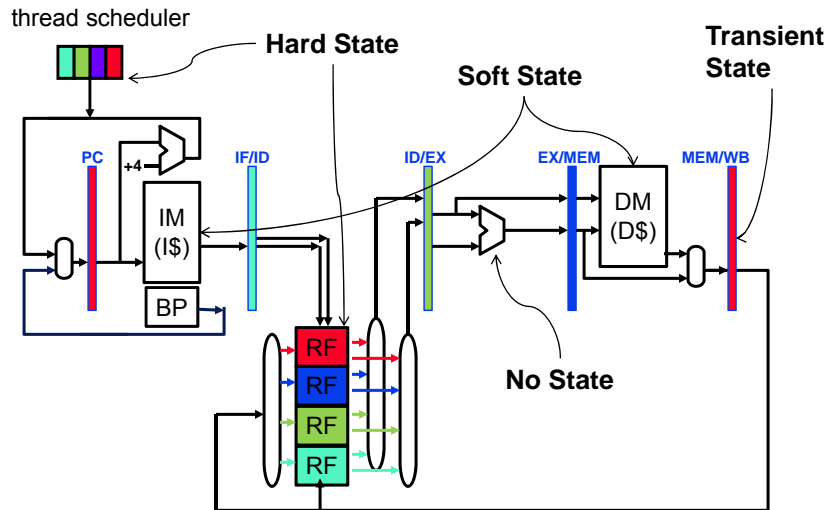
- Dynamically partitioned (like on a multi-programmed uni-processor)
 - TLBs need thread ids, caches/bpred table (BHT) don't
- Except **ordered "soft" state** (e.g., RAS) is replicated

❑ **Transient state**: pipeline latches

- Must be partitioned ... somehow

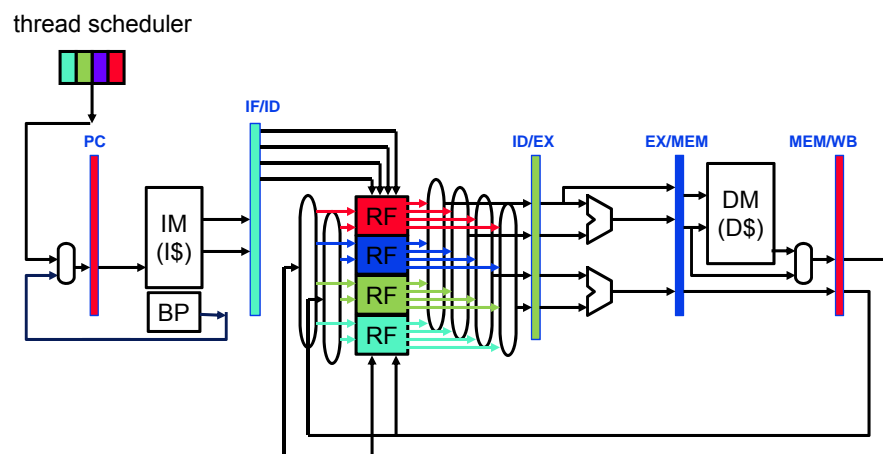
FGMT Datapath (Single Issue)

- ❑ What do we have to add to our datapath to support FGMT?

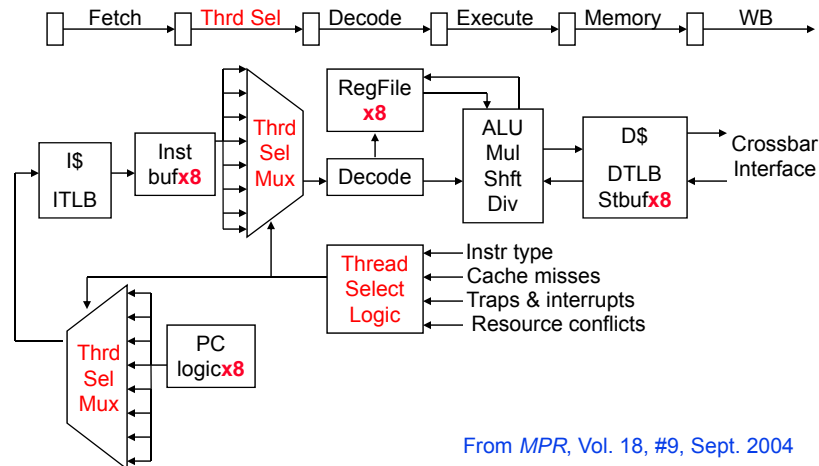


FGMT Datapath (2-Way Issue)

- ❑ What do we have to add to our datapath to support FGMT?



- ❑ Cores are simple (single-issue, 6 stage, no branch prediction), small, and power-efficient



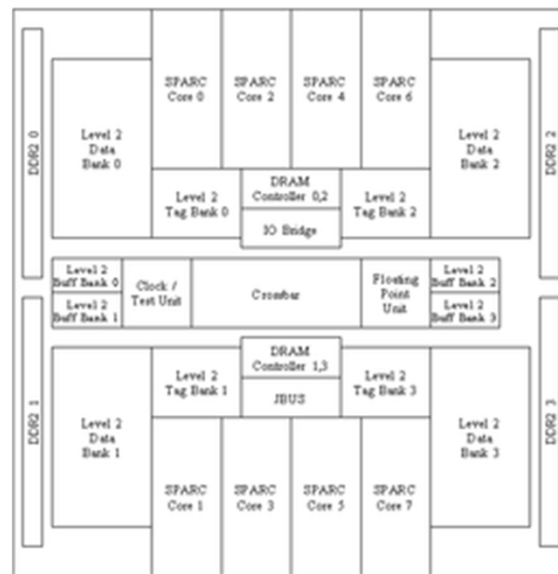
From *MPR*, Vol. 18, #9, Sept. 2004

CmpEn431 Chapter 4E.33

Irwin, PSU, 2015

Sun Niagara's Architecture

- 8 SPARC FGMT datapath cores



Niagra 1 / UltraSPARC T1 / OpenSPARC T1 - Die Micrograph Diagram (Seit 2006)

CmpEn431 Chapter 4E.34

Irwin, PSU, 2015

Reminders

❑ Next week

- Dynamic SuperScalar processors – P&H 4.10
 - out-of-order and SMT (aka hyperthreading)

❑ Reminders

- HW4 is due Oct 26th
- Quiz4 will close on Oct 23rd
- Second (and last) evening exam scheduled
 - Tuesday, **November 17**, 20:15 to 22:15, Location 22 Deike
 - Please let me know ASAP if you have a conflict !!