
CSE 431 Computer Architecture Fall 2015

Chapter 5B: Exploiting the Memory Hierarchy: Caches

Mary Jane Irwin (www.cse.psu.edu/~mji)

[Adapted from *Computer Organization and Design, 5th Edition*,
Patterson & Hennessy, © 2014, Morgan Kaufmann]

Reminders

□ Today

- Cache basics, improving cache performance – P&H, 5.3, 5.7

□ Next week

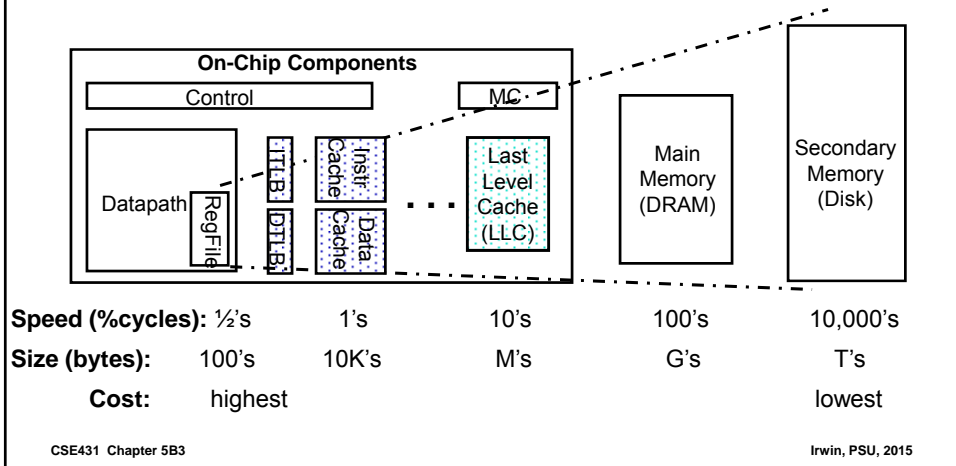
- Virtual memory hardware support (TLBs) – P&H 5.6-5.8
- VLIW datapaths – P&H 4.10

□ Reminders

- HW3 dropbox closes midnight Oct 1st
- Quiz 3 dropbox closes midnight Oct 4th
- First evening midterm exam scheduled
 - Tuesday, **October 6th**, 20:15 to 22:15, Location 22 Deike
 - No conflict exam !!

Review: A Typical Memory Hierarchy

- Take advantage of the **principle of locality** to present the user with as much memory as is available in the *cheapest* technology at the speed offered by the *fastest* technology



How is the Hierarchy Managed?

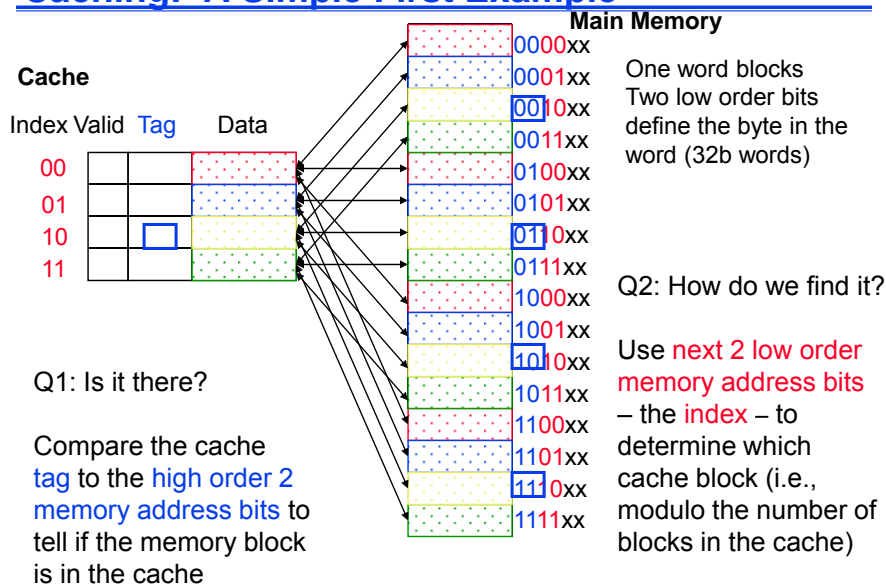
- registers ↔ memory
 - by compiler (programmer?)
- cache ↔ main memory
 - by the cache controller hardware
- main memory ↔ disks
 - by the operating system (virtual memory)
 - virtual to physical address mapping assisted by the hardware (TLB)
 - by the programmer (files)

Cache Basics

- ❑ Two questions to answer (in hardware):
 - Q1: How do we know if a data item is in the cache?
 - Q2: If it is, how do we find it?
- ❑ Direct mapped
 - Each memory block is mapped to exactly one block in the cache
 - lots of memory blocks must **share** a block in the cache
 - Address mapping (to answer Q2):

$$(\text{block address}) \bmod (\# \text{ of blocks in the cache})$$
 - Have a **tag** associated with each cache block that contains the address information (the upper portion of the address) required to identify the block (to answer Q1)

Caching: A Simple First Example




Direct Mapped Cache

- Consider the main memory word reference string

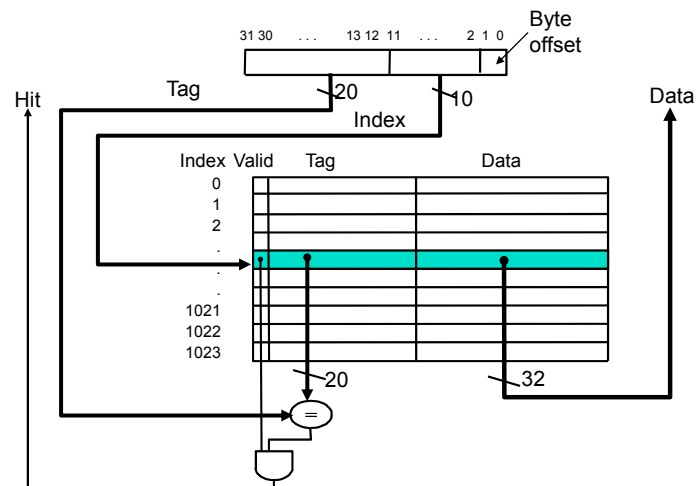
Start with an empty cache - all blocks initially marked as not valid

0 1 2 3 4 3 4 15



MIPS Direct Mapped 4KB Cache Example

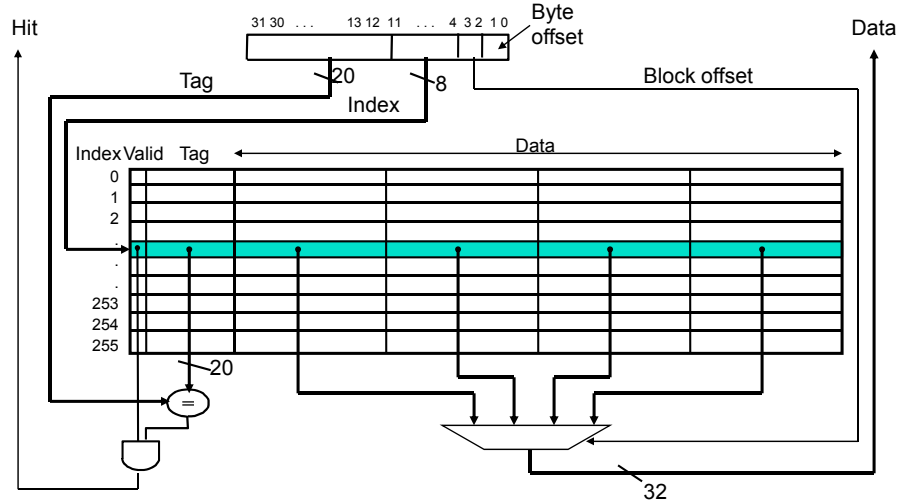
- One word blocks, cache size = 1K words (or 4KB)



What kind of locality are we taking advantage of?

Multiword Block Direct Mapped 4KB Cache

- Four words/block, cache size = 1K words



What kind of locality are we taking advantage of?

CSE431 Chapter 5B10


Irwin, PSU, 2015

Taking Advantage of Spatial Locality

- Let cache block hold more than one word

Start with an empty cache - all blocks initially marked as not valid

0 1 2 3 4 3 4 15




CSE431 Chapter 5B11

Irwin, PSU, 2015

Cache Field Sizes

- ❑ The number of bits in a cache includes both the storage for data and for the tags
 - 32-bit byte address
 - For a direct mapped cache with 2^n blocks, n bits are used for the index
 - For a block size of 2^m words (2^{m+2} bytes), m bits are used to address the word within the block and 2 bits are used to address the byte within the word
- ❑ What is the size of the tag field?
- ❑ The total number of bits in a direct-mapped cache is then
$$2^n \times (\text{block size} + \text{tag field size} + \text{valid field size})$$
- ❑ How many total bits are required for a direct mapped cache with 16KB of data and 4-word blocks assuming a 32-bit address?

Miss Rate vs Block Size vs Cache Size



- ❑ Miss rate goes up if the block size becomes a significant fraction of the cache size because the number of blocks that can be held in the same size cache is smaller (increasing **capacity** misses)


Handling Cache Hits

- ❑ Read hits (I\$ and D\$)
 - this is what we want!
- ❑ Write hits (D\$ only)
 - If we require the cache and memory to be **consistent**
 - always write the data into both the cache block and the next level in the memory hierarchy (**write-through**)
 - writes run at the speed of the next level in the memory hierarchy – so slow! – or can use a **write buffer** and stall only if the write buffer is full
 - If we allow cache and memory to be **inconsistent**
 - write the data only into the cache block (**write-back** the cache block to the next level in the memory hierarchy when that cache block is “evicted”)
 - need a **dirty** bit for each data cache block to tell if it needs to be written back to memory when it is evicted – can use a **write buffer** to help “buffer” write-backs of dirty blocks

Sources of Cache Misses

- ❑ **Compulsory** (cold start or process migration, first reference):
 - First access to a block, “cold” fact of life, not a whole lot you can do about it. If you are going to run “millions” of instructions, compulsory misses are insignificant
 - Solution: increase block size (increases miss penalty; very large blocks could increase miss rate)
- ❑ **Capacity**:
 - Cache cannot contain all blocks accessed by the program
 - Solution: increase cache size (may increase access time)
- ❑ **Conflict** (collision):
 - Multiple memory locations mapped to the same cache location
 - Solution 1: increase cache size
 - Solution 2: increase associativity (stay tuned) (may increase access time)

Miss Rates per Cache Miss Type



CSE431 Chapter 5B18

Irwin, PSU, 2015

Handling Cache Misses (Single Word Blocks)

- ❑ Read misses (I\$ and D\$)
 - **stall** the pipeline, fetch the block from the next level in the memory hierarchy, install it in the cache (which may involve having to evict a dirty block if using a write-back cache), and send the requested word to the core, then let the pipeline resume
- ❑ Write misses (D\$ only)
 1. **stall** the pipeline, fetch the block from next level in the memory hierarchy, install it in the cache (which may involve having to evict a dirty block if using a write-back cache), write the word from the core to the cache, then let the pipeline resume
 2. **Write allocate** – just write the word (and its tag) into the cache (which may involve having to evict a dirty block if using a write-back cache), no need to check for cache hit, no need to stall
 3. **No-write allocate** – skip the cache write (but must invalidate that cache block since it will now hold stale data) and just write the word to the write buffer (and eventually to the next memory level), no need to stall if the write buffer isn't full

CSE431 Chapter 5B19

Irwin, PSU, 2015

Multiword Block Considerations

❑ Read misses (I\$ and D\$)

- Processed the same as for single word blocks – a miss returns the entire block from memory
- Miss penalty grows as block size grows
 - **Early restart** – core resumes execution as soon as the requested word of the block is returned
 - **Requested word first** – requested word is transferred from the memory to the cache (and core) first
- **Nonblocking cache** – allows the core to continue to access the cache while the cache is handling an earlier miss

❑ Write misses (D\$ only)

- If using write allocate must *first* fetch the block from memory and then write the word to the block (or could end up with a “garbled” block in the cache (e.g., for 4 word blocks, a new tag, one word of data from the new block, and three words of data from the old block))

Measuring Cache Performance

❑ Assuming cache hit costs are included as part of the normal CPU execution cycle, then

$$\begin{aligned} \text{CPU time} &= IC \times CPI \times CC \\ &= IC \times \underbrace{(CPI_{\text{ideal}} + \text{Memory-stall cycles})}_{CPI_{\text{stall}}} \times CC \end{aligned}$$

❑ Memory-stall cycles come from cache misses (a sum of read-stalls and write-stalls)

$$\text{Read-stall cycles} = \text{reads/program} \times \text{read miss rate} \times \text{read miss penalty}$$

$$\begin{aligned} \text{Write-stall cycles} &= (\text{writes/program} \times \text{write miss rate} \\ &\quad \times \text{write miss penalty}) \\ &\quad + \text{write buffer stalls} \end{aligned}$$

❑ For write-through caches, we can simplify this to

$$\text{Memory-stall cycles} = \text{accesses/program} \times \text{miss rate} \times \text{miss penalty}$$

Impacts of Cache Performance

- ❑ Relative cache penalty increases as core performance improves (faster clock rate and/or lower CPI)
 - The memory speed is unlikely to improve as fast as core cycle time. When calculating CPI_{stall} , the cache miss penalty is measured in *core* clock cycles needed to handle a miss
 - The lower the CPI_{ideal} , the more pronounced the impact of stalls
- ❑ A core with a CPI_{ideal} of 2, a 100 cycle miss penalty, 36% load/store instr's, and 2% I\$ and 4% D\$ miss rates

$$\text{Memory-stall cycles} = 2\% \times 100 + 36\% \times 4\% \times 100 = 3.44$$

$$\text{So } CPI_{stalls} = 2 + 3.44 = \mathbf{5.44}$$

more than twice the CPI_{ideal} !

Impacts of Cache Performance, Con't

- ❑ Relative cache penalty increases as core performance improves (lower CPI)
- ❑ What if the CPI_{ideal} is reduced to 1? 0.5? 0.25?
- ❑ What if the D\$ miss rate went up 1%? 2%?
- ❑ What if the core clock rate is doubled (doubling the miss penalty)?

Average Memory Access Time (AMAT)

- ❑ A larger cache will have a longer access time. An increase in hit time will likely add another stage to the pipeline. At some point the increase in hit time for a larger cache will overcome the improvement in hit rate leading to a decrease in performance.
- ❑ Average Memory Access Time (AMAT) is the average to access memory considering both hits and misses

$$\text{AMAT} = \text{Time for a hit} + \text{Miss rate} \times \text{Miss penalty}$$

- ❑ What is the AMAT for a core with a 20 psec clock, a miss penalty of 50 clock cycles, a miss rate of 0.02 misses per instruction and a cache access time of 1 clock cycle?

$$\text{AMAT} = 1 + 0.02 \times 50 = 2 \text{ cycles}$$

Reducing Cache Miss Rates #1


1. Allow more flexible block placement
 - ❑ In a **direct mapped cache** a memory block maps to exactly one cache block
 - ❑ At the other extreme, could allow a memory block to be mapped to *any* cache block – **fully associative cache**
 - ❑ A compromise is to divide the cache into **sets** each of which consists of *n* “ways” (**n-way set associative**). A memory block maps to a unique set (specified by the index field) and can be placed in any way of that set (so there are *n* choices)

$$(\text{block address}) \bmod (\# \text{ sets in the cache})$$

Another Reference String Mapping

□ Consider the main memory word reference string

Start with an empty cache - all blocks initially marked as not valid



• 8 requests, 8 misses

□ Ping pong effect due to **conflict** misses - two memory locations that map into the same cache block

Set Associative Cache Example

Cache

Way	Set	V	Tag	Data
0	0			
	1			
1	0			
	1			

Q1: Is it there?

Compare *all* the cache tags in the set to the **high order 3 memory address bits** to tell if the memory block is in the cache

Main Memory

000	0xx
000	1xx
001	0xx
001	1xx
010	0xx
010	1xx
011	0xx
011	1xx
100	0xx
100	1xx
101	0xx
101	1xx
110	0xx
110	1xx
111	0xx
111	1xx

One word blocks
Two low order bits
define the byte in the
word (32b words)

Q2: How do we find it?


Use **next 1 low order memory address bit** to determine which cache set (i.e., modulo the number of sets in the cache)

Another Reference String Mapping

- Consider the main memory word reference string

Start with an empty cache - all blocks initially marked as not valid

0 4 0 4 0 4 0 4

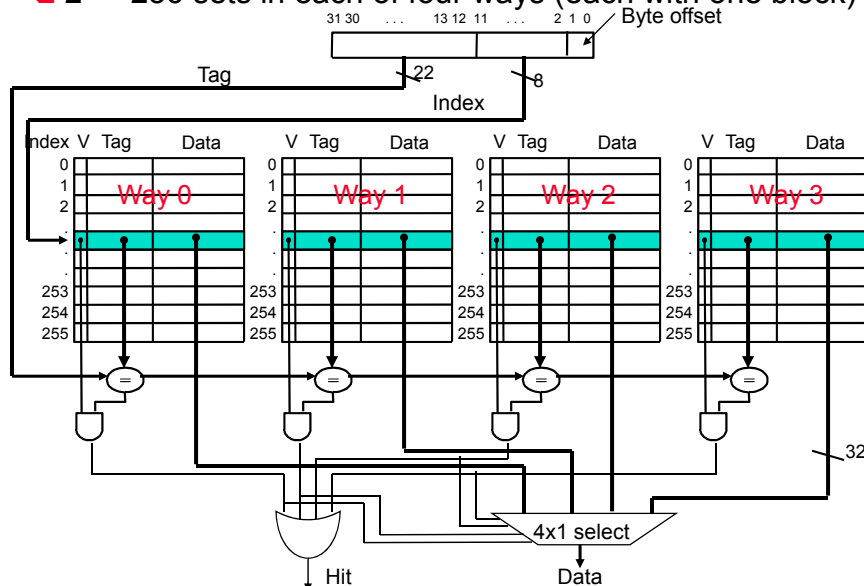


- 8 requests, 2 misses

- Solves the ping pong effect in a direct mapped cache due to **conflict** misses since now two memory locations that map into the same cache set can co-exist!

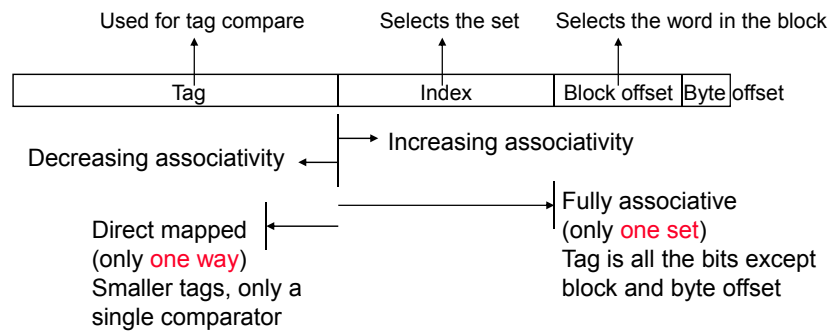
Four-Way Set Associative 4KB Cache

- $2^8 = 256$ sets in each of four ways (each with one block)



Range of Set Associative Caches

- For a **fixed size** cache, each increase by a factor of two in associativity doubles the number of ways (i.e., doubles the number of blocks per set) and halves the number of sets – decreases the size of the index by 1 bit and increases the size of the tag by 1 bit




Costs of Set Associative Caches

- When a miss occurs, which way's block do we pick for replacement?
 - Least Recently Used (LRU): the block replaced is the one that has been unused for the longest time
 - Must have hardware to keep track of when each way's block was used relative to the other blocks in the set
 - For 2-way set associative, takes **one bit per way** → set the bit when a block is referenced (and reset the other way's bit)
- N-way set associative cache additional costs
 - N comparators (delay and area) – one per way
 - MUX delay (way selection) before data is available
 - Data available **after** way selection (and Hit/Miss decision). In a direct mapped cache, the cache block is available **before** the Hit/Miss decision
 - So its not possible to just assume a hit and continue and recover later if it was a miss

Benefits of Set Associative Caches

- ❑ The choice of direct mapped or set associative depends on the cost of a miss versus the cost of implementation



- ❑ Largest gains are in going from direct mapped to 2-way (20%+ reduction in miss rate)

Reducing Cache Miss Rates #2

2. Use multiple levels of caches

- ❑ With advancing technology have more than enough room on the die for bigger L1 caches *and* for a second level of caches – normally a **unified** L2 cache (i.e., it holds both instructions and data) and even a **unified** L3 cache
- ❑ For our example, CPI_{ideal} of 2, 100 cycle miss penalty (to main memory) and a 25 cycle miss penalty (to UL2\$), 36% load/stores, a 2% (4%) L1I\$ (L1D\$) miss rate, add a 0.5% UL2\$ miss rate

$$CPI_{stalls} = 2 + .02 \times 25 + .36 \times .04 \times 25 + .005 \times 100 + .36 \times .005 \times 100 = 3.54$$

(as compared to 5.44 with no L2\$)

Multilevel Cache Design Considerations

- ❑ Design considerations for L1 and L2 caches are very different
 - Primary cache should focus on **minimizing hit time** in support of a shorter core clock cycle
 - Smaller with smaller block sizes
 - Secondary cache(s) should focus on **reducing miss rate** to reduce the penalty of long main memory access times
 - Larger with larger block sizes
 - Higher levels of associativity
- ❑ The miss penalty of the L1 cache is significantly reduced by the presence of an L2 cache – so it can be smaller (i.e., faster) but have a higher miss rate
- ❑ For the L2 cache, hit time is less important than miss rate
 - The L2\$ hit time determines L1's miss penalty
 - L2\$ local miss rate \gg than the global miss rate

Split vs Unified Caches


- ❑ Split L1I\$ and L1D\$: instr's and data in different caches at L1
 - To minimize structural hazards and t_{hit}
 - So low capacity/associativity (to reduce t_{hit})
 - So small to medium block size (to reduce conflict misses)
 - To optimize L1I\$ for wide output (superscalar) and no writes
- ❑ Unified L2, L3, ...: instr's and data together in one cache
 - To minimize $\%_{miss}$ (t_{hit} is less important due to (hopefully) infrequent accesses)
 - So high capacity/associativity/block size (to reduce $\%_{miss}$)
 - Fewer capacity misses: unused instr capacity can be used for data
 - More conflict misses: instr / data conflicts (smaller effect in large caches)
 - Instr / data structural hazards are rare (would take a simultaneous L1I\$ and L1D\$ miss)

Comparing Cache Memory Architectures


Characteristic	ARM Cortex-A8	Intel Nehalem
L1 cache organization	Split instruction and data caches	Split instruction and data caches
L1 cache size	32 KiB each for instructions/data	32 KiB each for instructions/data per core
L1 cache associativity	4-way (I), 4-way (D) set associative	4-way (I), 8-way (D) set associative
L1 replacement	Random	Approximated LRU
L1 block size	64 bytes	64 bytes
L1 write policy	Write-back, Write-allocate(?)	Write-back, No-write-allocate
L1 hit time (load-use)	1 clock cycle	4 clock cycles, pipelined
L2 cache organization	Unified (instruction and data)	Unified (instruction and data) per core
L2 cache size	128 KiB to 1 MiB	256 KiB (0.25 MiB)
L2 cache associativity	8-way set associative	8-way set associative
L2 replacement	Approximated LRU	Approximated LRU
L2 block size	64 bytes	64 bytes
L2 write policy	Write-back, Write-allocate	Write-back, Write-allocate
L2 hit time	10 clock cycles	10 clock cycles
L3 cache organization	Unified (instruction and data)	Unified (instruction and data)
L3 cache size	8 MiB, shared	8 MiB, shared
L3 cache associativity	16-way set associative	16-way set associative
L3 replacement	Approximated LRU	Approximated LRU
L3 block size	64 bytes	64 bytes
L3 write policy	Write-back, Write-allocate	Write-back, Write-allocate
L3 hit time	35 clock cycles	35 clock cycles

Intel i7
(Haswell, 4core, 2wayHT)
 3.4GHz, 84W
 L1D\$ 4x32KB, 8-way
 L1I\$ 4x32KB, 8-way
 L2 4x256KB, 8-way
 L3 1x8MB, 16-way
 Cycles: L1:4, L2:11-16, L3:30-55
 1600MHz DDR3

ARM Cortex-A8 Data Cache Miss Rates



Intel i7 Data Cache Miss Rates



CSE431 Chapter 5B39

Irwin, PSU, 2015

Reducing Cache Miss Rates #3

3. Hardware prefetching

- Fetch blocks into the cache proactively (speculatively)
- Key is to anticipate the upcoming miss addresses accurately
- Relies on having unused memory bandwidth available

□ A simple case is to use **next block** prefetching

- Miss on address $X \rightarrow$ anticipate next reference miss on $X + \text{block-size}$
- Works well for instr's (sequential execution) and for arrays of data

□ Need to initiate prefetches sufficiently in advance

□ If prefetch instr/data that is not going to be used then have polluted the cache with unnecessary data (possibly evicting useful data)


CSE431 Chapter 5B40

Irwin, PSU, 2015

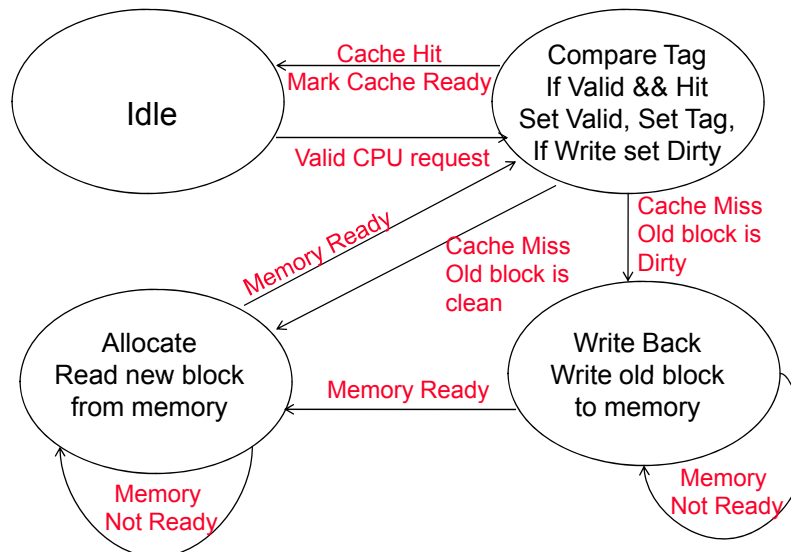
FSM Cache Controller

□ Key characteristics for a simple L1 cache

- Direct mapped
- Write-back using write-allocate
- Block size of 4 32-bit words (so 16B); Cache size of 16KB (so 1024 blocks)
- 18-bit tags, 10-bit index, 2-bit block offset, 2-bit byte offset, dirty bit, valid bit, LRU bits (if set associative)



Four State Cache Controller



Cache Coherence Issues – More Details to Come

- ❑ Need a cache controller to also ensure cache coherence the most popular of which is **snooping**
 - The cache controller monitors (snoops) on the broadcast medium (e.g., bus) with duplicate address tag hardware (so it can determine if it is valid)
- ❑ **Will cover this important material later in the semester**
 - Write access: If multiple copies of an item exist, the first write invalidates all other copies
- ❑ If two cores attempt to write the same data at the same time, one of them wins the race causing the other core's copy to be invalidated. For the other core to complete, it must obtain a new copy of the data which must now contain the updated value – thus enforcing **write serialization**


Impact of Code Transformations on Cache Performance

- ❑ Code transformations change data access pattern, influencing cache hits and misses
- ❑ A large body of compiler transformations designed to maximize cache performance
 - **Data Reuse => Data Locality => Cache Performance**
- ❑ Examples
 - Loop interchange
 - Iteration space tiling (aka code blocking)
 - Loop fusion
 - Statement scheduling

Loop Fusion

- ❑ Merges two adjacent countable loops into a single loop
- ❑ Reduces the cost of test and branch code
- ❑ Fusing loops that refer to the same data enhances temporal locality
- ❑ One potential drawback is that larger loop body may reduce instruction locality when the instruction cache is very small

```
for i = 1, N
  A(i) = B(i) + C(i)
for i = 1, N
  D(i) = E(i) + B(i)
```




```
for i = 1, N
  A(i) = B(i) + C(i)
  D(i) = E(i) + B(i)
```

Loop Interchange

- ❑ Changes the direction of array traversing by swapping two loops.
- ❑ The goal is to align data access direction with the memory layout order.

```
for i = 1, N
  for j = 1, N
    ... A(j,i) ...
```



```
for j = 1, N
  for i = 1, N
    ... A(j,i) ...
```

assuming ROW-MAJOR memory layout

- ❑ What type of locality do we exploit?

Summary: Improving Cache Performance

0. Reduce the time to hit in the cache

- smaller cache
- direct mapped cache
- smaller blocks
- for writes
 - no write allocate – no “hit” on cache, just write to write buffer
 - write allocate – to avoid two cycles (first check for hit, then write) pipeline writes via a delayed write buffer to cache

1. Reduce the miss rate

- bigger cache
- more flexible placement (increase associativity)
- larger blocks (16 to 64 bytes typical)
- victim cache – small buffer holding most recently discarded blocks

Summary: Improving Cache Performance


2. Reduce the miss penalty

- smaller blocks
- use a write buffer to hold dirty blocks being replaced so don't have to wait for the write to complete before reading
- check write buffer (and/or victim cache) on read miss – may get lucky
- for large blocks fetch critical word first
- use multiple cache levels – L2 cache not tied to CPU clock rate
- faster backing store/improved memory bandwidth
 - wider buses
 - memory interleaving, DDR SDRAMs

Summary: The Cache Design Space


❑ Several interacting dimensions

- cache size
- block size
- associativity
- replacement policy
- write-through vs write-back
- write allocation



❑ The optimal choice is a compromise

- depends on access characteristics
 - workload
 - use (I-cache, D-cache, TLB)
- depends on technology / cost



❑ Simplicity often wins

Reminders

❑ Next week

- Virtual memory hardware support (TLBs) – P&H 5.6-5.8
- VLIW datapaths – P&H 4.10

❑ Reminders

- HW3 dropbox closes midnight Oct 1st
- Quiz 3 dropbox closes midnight Oct 4th
- First evening midterm exam scheduled
 - Tuesday, **October 6th**, 20:15 to 22:15, Location 22 Deike
 - No conflict exam !!